

SUPPLEMENT for Submission 6580:
Parallel Greedy Best-First Search with a Bound on the Number of Expansions
Relative to Sequential Search

Primary Keywords: *None*

Abstract

Supplementary material for submission.

S.1 K-Parallel GBFS (KPGBFS)

Algorithm S.1: K-Parallel GBFS (KPGBFS)

```

1:  $Open \leftarrow \{s_{init}\}, Closed \leftarrow \{s_{init}\}; \forall i, s_i \leftarrow NULL$ 
2: for  $i \leftarrow 0, \dots, k-1$  in parallel do
3:   loop
4:     while  $s_i = NULL$  do
5:        $lock(Open)$ 
6:       if  $\forall j, s_j = NULL$  then
7:         if  $Open = \emptyset$  then  $unlock(Open); return NULL$ 
8:       else
9:          $s_i \leftarrow top(Open); Open \leftarrow Open \setminus \{s_i\}$ 
10:       $unlock(Open)$ 
11:      if  $s_i \in S_{goal}$  then return  $s_i$ 
12:       $lock(Open), lock(Closed)$ 
13:      for  $s'_i \in succ(s_i)$  do
14:        if  $s'_i \notin Closed$  then
15:           $Closed \leftarrow Closed \cup \{s'_i\}$ 
16:           $Open \leftarrow Open \cup \{s'_i\};$ 
17:       $unlock(Open), unlock(Closed)$ 
18:       $s_i \leftarrow NULL$ 

```

$\triangleright k$ is the number of the threads

S.2 PUHF Pathology with Consistent Heuristics

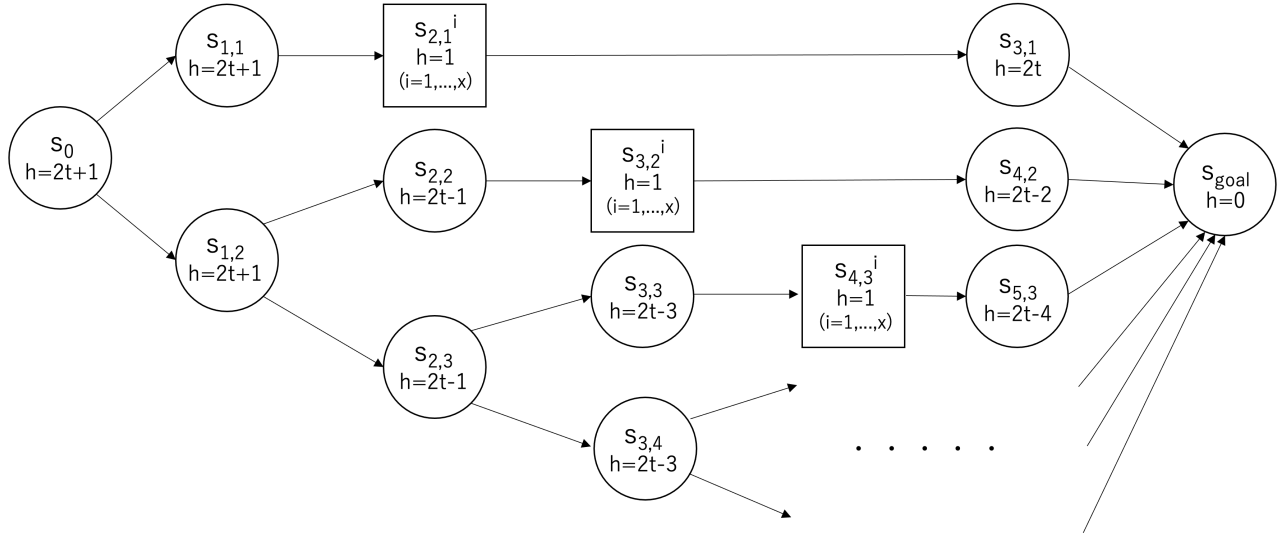


Figure 1: PUHF Pathological Search Behavior Example 2 (same as Figure 2 in the main paper)

PUHF is pathological relative to GBFS even if the heuristic is consistent. For example, consider an instance of the search space with the structure shown in Figure 1, where all edge costs are 1. If we replace all h -values by h' , such that $h' = h/c$, where c is a sufficiently large constant such that $h'/c < 1$, then h' is consistent, but PUHF is pathological on this search space (by the same argument as the main paper).

S.3 OBAT with SGE (OBAT_S)

Algorithm S.2: OBAT with SEE (OBAT_S), lock/unlock operations for *Being_Expanded* omitted for space

```

1:  $Open \leftarrow \{s_{init}\}, Closed \leftarrow \{s_{init}\}; \forall i, s_i \leftarrow NULL$ 
2: for  $i \leftarrow 0, \dots, k-1$  in parallel do
3:   loop
4:     lock(Unevaluated)
5:     if Unevaluated  $\neq \emptyset$  then
6:        $s_i \leftarrow top(Unevaluated)$ 
7:        $Unevaluated \leftarrow Unevaluated \setminus \{s_i\}$ 
8:       unlock(Unevaluated)
9:       evaluate( $s_i$ ) ▷ using a hashtable to prevent reevaluation of states
10:       $s \leftarrow parent(s_i)$ 
11:      if all succ( $s$ ) have been evaluated then
12:        if  $h(s) \leq \min_{s' \in succ(s)} h(s')$  then ▷  $s$  is an  $a$ -state
13:          lock(Open), lock(Closed)
14:          for  $s' \in succ(s)$  do
15:            if  $s' \notin Closed$  then
16:               $Closed \leftarrow Closed \cup \{s'\}$ 
17:               $Open \leftarrow Open \cup \{s'\}$ 
18:            unlock(Open), unlock(Closed)
19:          else ▷  $s$  is a  $b$ -state
20:            lock(Deferred)
21:             $Deferred \leftarrow Deferred \cup \{s\}$ 
22:            unlock(Deferred)
23:             $Being\_Expanded \leftarrow Being\_Expanded \setminus \{s\}$ 
24:          else
25:            unlock(Unevaluated)
26:            lock(Open), lock(Deferred)
27:            if Open =  $\emptyset$  and Deferred =  $\emptyset$  then
28:              unlock(Open), unlock(Deferred)
29:              if Being_Expanded =  $\emptyset$  then return NULL
30:            else
31:              if  $h(top(Deferred)) \leq h(top(Open))$  then
32:                if  $h(top(Deferred)) \leq \min_{s \in Being\_Expanded} h(s)$  then
33:                   $s_i \leftarrow top(Deferred); Deferred \leftarrow Deferred \setminus \{s_i\}$ 
34:                   $Being\_Expanded \leftarrow Being\_Expanded \cup \{s_i\}$ 
35:                  lock(Closed)
36:                  for  $s'_i \in succ(s_i)$  do
37:                    if  $s'_i \notin Closed$  then
38:                       $Closed \leftarrow Closed \cup \{s'_i\}$ 
39:                       $Open \leftarrow Open \cup \{s'_i\}$ 
40:                    unlock(Closed)
41:                   $Being\_Expanded \leftarrow Being\_Expanded \setminus \{s_i\}$ 
42:                unlock(Open), unlock(Deferred)
43:              else
44:                if  $h(top(Open)) \leq \min_{s \in Being\_Expanded} h(s)$  then
45:                   $s_i \leftarrow top(Open); Open \leftarrow Open \setminus \{s_i\}$ 
46:                   $Being\_Expanded \leftarrow Being\_Expanded \cup \{s_i\}$ 
47:                  unlock(Open), unlock(Deferred)
48:                  if  $s_i \in S_{goal}$  then return Path( $s_i$ )
49:                  generate(succ( $s_i$ ))
50:                  lock(Unevaluated)
51:                  for  $s'_i \in succ(s_i)$  do
52:                     $Unevaluated \leftarrow Unevaluated \cup \{s'_i\}$ 
53:                     $parent(s'_i) \leftarrow s_i$ 
54:                  unlock(Unevaluated)
55:                  if succ( $s_i$ ) =  $\emptyset$  then  $Being\_Expanded \leftarrow Being\_Expanded \setminus \{s_i\}$ 
56:                else
57:                  unlock(Open), unlock(Deferred)
58:       $s_i \leftarrow NULL$ 

```

Algorithm S.2 shows OBAT_S (OBAT with SGE).

Although the pseudocode may appear more complex than OBAT, the additional code required for OBAT_S is a straightforward implementation of SGE, very similar to the previous application of SGE to PUHF and KPGBFS (Shimoda and Fukunaga 2024).

Most of the new pseudocode relative to OBAT are in the `if`-block in lines 4–23, which handles the parallel evaluation of states in *Unevaluated* by available threads.

In OBAT, when a state s leaves *Open*, the thread which removed s generates $\text{succ}(s)$, the successors of s , and evaluates the h -values of $\text{succ}(s)$.

In OBAT_S , when a state s leaves *Open*, the thread which removed s generates $\text{succ}(s)$, but instead of evaluating $\text{succ}(s)$, inserts $\text{succ}(s)$ into *Unevaluated* (lines 50-54). The unevaluated states in *Unevaluated* are evaluated and processed as soon as threads become available, having higher priority than removing states from *Deferred* and *Open*.

Being_Expanded is an auxiliary data structure which contains the set of states which are currently being expanded according to Definition 10.

S.4 Correctness of OBAT_S

Applying SGE to OBAT, KPGBFS, and PUHF can change the search behavior. In the case of KPGBFS, where state expansion is unconstrained and the state with best h -value in the shared *OPEN* is always expanded, there is no known constraint on the set of states which can be expanded, so adding SGE has no relevant theoretical implication.

In the case of PUHF and its variants, including PUHF3, it was argued in (Shimoda and Fukunaga 2024) that although SGE can change the order in which states are expanded, SGE can not cause PUHF to violate the state expansion constraint, i.e., PUHF with SGE can not expand a state which is not in the BTS. This is because PUHF explicitly checks the state expansion constraint before expanding a state, so it is not possible for PUHF with SGE to expand a state which does not satisfy the expansion constraint.

Now, let us consider OBAT with SGE. It is clear that like PUHF with SGE, adding SGE to OBAT can not cause the search to violate the two main constraints (1) only states in the BTS are expanded, and (2) only 1 bench is explored at a time, because these constraints are enforced by the `if`-statement conditions in Algorithm S.2, lines 31, 32, 44, and these are the same conditions as in OBAT (Algorithm 1, lines 10, 11, 21).

The proofs for Theorem 2 and Lemma 1 depends only on the expansion constraints and priority ordering scheme for removing states from *Open* and *Deferred*, which are the same in OBAT and OBAT_S , so Theorem 2 and Lemma 1 hold for OBAT_S .

It remains to show Lemma 2 (and therefore Theorem 4) holds for OBAT_S .

We need to show that the number of states remaining in *Deferred* after OBAT_S terminates does not exceed $k|p|$, the bound given by Theorem 3.

As OBAT_S differs from OBAT in that SGE separates successor generation and evaluation such that successors are not necessarily immediately evaluated by the same thread which generated them, we define the following:

Definition 10. We say that a state s is *being expanded* if either:

- s has been removed from *Open*, at least one member of $\text{succ}(s)$ has not yet been evaluated. or
- s has been removed from *Deferred*, at least one member of $\text{succ}(s)$ has not yet been inserted in *Open*.

A thread might remove s from *Open*, insert $\text{succ}(s)$ in *Unevaluated*, and move on to remove another state from *Open* or *Deferred* while s is still “being expanded” by the definition above. However, we show that there are never more than k states being expanded during OBAT_S search.

Lemma 3. The maximum number of states being expanded by OBAT_S is at most k .

Proof. At any given time, let n be the number of threads which are expanding a state from *Deferred* (Alg S.2 Lines 33-41) or generating its successors (Alg S.2, Lines 45-49), and let m be the number of states whose successors have been inserted into *Open*, but at least one member of successors has not yet been evaluated.

$n + m$ is the number of states being expanded by OBAT_S at any given time. We show that $n + m \leq k$ by induction.

First, the condition trivially holds at the start of search ($n = 0, m = 0$). Next, we perform induction on n and m . Assume the condition $n + m \leq k$ holds. Consider when n increases: *Unevaluated* is prioritized over *Open* and *Deferred*, so *Unevaluated* is empty. Let l be the number of threads currently evaluating states. Clearly, $m \leq l$. Since $n + l \leq k$, $n + m \leq k$. Consider when m increases: When some thread i inserts a successor state $u \in \text{succ}(s_i)$ into *Unevaluated*, m increases by 1. This also means that thread i has finished generating the $\text{succ}(s_i)$ so n decreases by 1, and $n + m$ is unchanged, so $n + m \leq k$. This completes the induction. \square

This bound on the number of states being expanded allows us to obtain the equivalent of Lemma 2 for OBAT_S .

Lemma 4. In OBAT_S , for every h -value, *Deferred* contains at most k states, where k is the number of threads.

Proof. Assume that at some point during search, $k + 1$ states which have the same h -value are inserted in *Deferred*. Let n_1, n_2, \dots, n_{k+1} denote the order in which these states were inserted in *Deferred*. 60

After n_1 is inserted in *Deferred*, it is not possible to remove n_i ($2 \leq i \leq k + 1$) from *Open*. This is because if n_i ($2 \leq i \leq k + 1$) is removed from *Open* after n_1 is inserted in *Deferred*, then by the time n_i is removed from *Open*, n_1 must already have been removed from *Deferred*, because among states with the same h -value, removal from *Deferred* is prioritized over *Open*. This means n_1 and n_i could not both have been simultaneously in *Deferred*, a contradiction. 65

Thus, immediately before n_1 is inserted in *Deferred*, n_i ($2 \leq i \leq k + 1$) must have been removed from *Open*, and not yet inserted into *Deferred*. i.e., these states are being expanded. Thus, $k + 1$ states (n_1, \dots, n_{k+1}) must be simultaneously expanded, but at most k states are being expanded, a contradiction. \square

From Lemmas 1, 3, and 4, Theorem 3 holds for OBAT_S , so Theorem 4 also holds for OBAT_S . 70

S.5 Explanation of Benchmark Selection

We compared GBFS, OBAT, OBAT_S , KPGBFS, KPGBFS_S (KPGBFS with SGE), PUHF3, PUHF3_S (PUHF3 with SGE) using a set of instances based on the Autoscale-21.11 satisficing benchmark set (42 STRIPS domains, 30 instances/domain, 1260 total instances) (Torrallba, Seipp, and Sievers 2021). The Autoscale benchmarks are an improved benchmark suite based on the IPC classical planning benchmarks which were designed to compare the performance of different solvers, as advances in solvers sometimes made performance comparisons among modern solvers difficult using the classic IPC competition benchmark instances. 75

However, even on the Autoscale-21.11 benchmark set, there were several domains which were too easy – all methods solved all instances, rendering these instances useless for the purpose of comparing coverage among the methods.

Therefore, for domains where (1) all methods solved all instances for $k = 4$, $k = 8$, and $k = 16$ threads, and (2) a parameterized instance generator for the domain is available in the Autoscale repository, we replaced the Autoscale-21.11 instances with more difficult instances generated using the same Autoscale instance generator. Specifically, the domains where criteria (1) and (2) above applied were gripper, and miconic. 80

- gripper: The Autoscale-21.11 set used values of parameter n from 20 to 165, in increments of 5 (30 instances). Our instances used n from 175 to 465, in increments of 10 (30 instances). 85

- miconic:

The Autoscale miconic instances are generated using two parameters, *passengers* and *floors*. In the Autoscale-21.11 set (30 instances), $19 \leq \text{passengers} \leq 155$, where *passengers* monotonically increased by 4 or 5 between consecutive instances, and $11 \leq \text{floors} \leq 124$, monotonically increasing in increments of 3 or 4.

Similarly, our replacement instances were generated using the following procedure, which calls the Autoscale `generate_instance` function for miconic: 90

```
passengers=155, floors=124
for i in range(0,90):
    if(i%4==0):
        passengers+=4
    else:
        passengers+=5
    floors+=4
    if(i%3==2):
        generate_instance(passengers,floors)
```

100

S.6 Coverage Per Domain (full table)

| #threads | 1 thread | 4 threads | | | | | | 8 threads | | | | | | 16 threads | | | | | |
|-------------------------|----------|-----------|---------------------|-----------|--------------------|-----------|-------------------|-----------|---------------------|-----------|--------------------|-----------|-------------------|------------|---------------------|-----------|--------------------|-----------|-------------------|
| | GBFS | KPGBFS | KPGBFS _S | PUHF3 | PUHF3 _S | OBAT | OBAT _S | KPGBFS | KPGBFS _S | PUHF3 | PUHF3 _S | OBAT | OBAT _S | KPGBFS | KPGBFS _S | PUHF3 | PUHF3 _S | OBAT | OBAT _S |
| agricola | 26 | 30 | 30 | 30 | 28 | 28 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 |
| airport | 18 | 19 | 19 | 18 | 18 | 19 | 19 | 18 | 19 | 17 | 19 | 16 | 19 | 17 | 18 | 17 | 19 | 15 | 18 |
| barman | 2 | 3 | 3 | 3 | 3 | 3 | 3 | 4 | 4 | 4 | 4 | 3 | 4 | 4 | 4 | 4 | 4 | 4 | 5 |
| blocksworld | 6 | 6 | 6 | 7 | 6 | 6 | 6 | 7 | 6 | 7 | 7 | 8 | 7 | 7 | 7 | 7 | 8 | 6 | 7 |
| childsnaek | 4 | 3 | 5 | 3 | 4 | 5 | 3 | 5 | 4 | 5 | 4 | 3 | 3 | 5 | 3 | 5 | 4 | 4 | 4 |
| data-network | 4 | 6 | 5 | 5 | 5 | 5 | 5 | 7 | 5 | 5 | 5 | 4 | 5 | 7 | 6 | 6 | 4 | 5 | 5 |
| depots | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 5 | 5 | 4 | 4 | 4 | 4 |
| driverlog | 4 | 8 | 7 | 7 | 7 | 7 | 8 | 7 | 7 | 7 | 7 | 6 | 8 | 7 | 8 | 8 | 8 | 7 | 9 |
| elevators | 11 | 15 | 14 | 16 | 16 | 15 | 14 | 17 | 17 | 17 | 17 | 16 | 18 | 18 | 19 | 20 | 20 | 20 | 19 |
| floortile | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| freecell | 20 | 21 | 21 | 24 | 21 | 23 | 22 | 22 | 23 | 21 | 23 | 22 | 23 | 24 | 23 | 21 | 22 | 22 | 25 |
| ged | 7 | 8 | 8 | 7 | 9 | 9 | 9 | 9 | 8 | 8 | 9 | 10 | 9 | 9 | 9 | 7 | 8 | 10 | 8 |
| grid | 5 | 6 | 5 | 5 | 6 | 5 | 6 | 5 | 5 | 6 | 6 | 6 | 6 | 5 | 6 | 6 | 6 | 6 | 6 |
| gripper (replaced) | 2 | 4 | 10 | 4 | 10 | 10 | 10 | 7 | 15 | 7 | 15 | 14 | 15 | 7 | 21 | 7 | 20 | 19 | 20 |
| hiking | 4 | 7 | 7 | 7 | 6 | 6 | 6 | 10 | 7 | 8 | 7 | 6 | 7 | 14 | 9 | 11 | 8 | 10 | 9 |
| logistics | 4 | 5 | 5 | 5 | 6 | 5 | 5 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 5 | 5 | 6 | 6 |
| miconic (replaced) | 10 | 9 | 13 | 10 | 13 | 10 | 15 | 9 | 15 | 9 | 15 | 10 | 18 | 9 | 15 | 9 | 15 | 10 | 21 |
| mprime | 4 | 5 | 6 | 4 | 4 | 4 | 5 | 5 | 6 | 6 | 5 | 5 | 5 | 6 | 6 | 6 | 6 | 5 | 5 |
| nomystery | 6 | 11 | 8 | 9 | 9 | 5 | 7 | 14 | 9 | 10 | 8 | 4 | 7 | 14 | 11 | 11 | 11 | 3 | 6 |
| openstacks | 7 | 9 | 9 | 9 | 9 | 15 | 17 | 11 | 9 | 10 | 10 | 15 | 20 | 12 | 11 | 11 | 11 | 15 | 22 |
| organic-synthesis-split | 9 | 16 | 17 | 16 | 16 | 12 | 16 | 17 | 17 | 16 | 17 | 13 | 17 | 17 | 17 | 16 | 16 | 13 | 16 |
| parcprinter | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 |
| parking | 6 | 8 | 7 | 6 | 6 | 7 | 6 | 7 | 8 | 8 | 8 | 8 | 8 | 9 | 10 | 9 | 9 | 9 | 10 |
| pathways | 11 | 12 | 12 | 12 | 13 | 11 | 12 | 11 | 15 | 12 | 13 | 10 | 9 | 15 | 12 | 10 | 12 | 11 | 12 |
| pegasol | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 30 | 29 | 30 | 30 | 30 | 30 | 30 | 30 |
| pipesworld-notankage | 8 | 8 | 7 | 9 | 8 | 9 | 8 | 6 | 8 | 10 | 9 | 10 | 7 | 11 | 9 | 9 | 9 | 9 | 11 |
| pipesworld-tankage | 9 | 11 | 12 | 12 | 12 | 10 | 10 | 10 | 12 | 12 | 12 | 10 | 10 | 12 | 12 | 11 | 11 | 12 | 12 |
| rovers | 26 | 26 | 26 | 26 | 26 | 26 | 26 | 27 | 25 | 26 | 25 | 26 | 26 | 27 | 27 | 26 | 26 | 26 | 26 |
| satellite | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 |
| scanalyzer | 7 | 9 | 9 | 9 | 9 | 9 | 9 | 10 | 10 | 10 | 10 | 10 | 10 | 11 | 11 | 11 | 11 | 11 | 11 |
| snake | 7 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 10 | 10 | 9 | 9 | 9 | 9 |
| sokoban | 16 | 18 | 19 | 17 | 18 | 18 | 18 | 21 | 22 | 19 | 18 | 19 | 19 | 23 | 22 | 21 | 21 | 19 | 19 |
| storage | 3 | 2 | 3 | 3 | 3 | 3 | 4 | 3 | 3 | 4 | 3 | 4 | 4 | 2 | 4 | 4 | 4 | 3 | 4 |
| termes | 12 | 15 | 16 | 16 | 16 | 14 | 16 | 18 | 18 | 16 | 17 | 15 | 16 | 18 | 18 | 16 | 18 | 15 | 16 |
| tetris | 8 | 9 | 9 | 9 | 9 | 10 | 9 | 11 | 11 | 11 | 11 | 11 | 12 | 13 | 12 | 13 | 12 | 13 | 13 |
| thoughtful | 14 | 18 | 18 | 18 | 16 | 16 | 15 | 18 | 17 | 16 | 17 | 17 | 17 | 21 | 18 | 19 | 16 | 16 | 14 |
| tidybot | 12 | 11 | 12 | 12 | 12 | 12 | 14 | 12 | 11 | 11 | 12 | 13 | 15 | 14 | 13 | 12 | 11 | 14 | 13 |
| tpp | 8 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 10 | 10 | 10 | 11 | 9 | 10 | 11 | 12 |
| transport | 5 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 8 | 7 | 7 | 7 | 7 | 8 | 7 | 8 | 7 | 7 | 7 |
| visitall | 14 | 15 | 15 | 12 | 15 | 12 | 16 | 15 | 18 | 14 | 16 | 15 | 16 | 20 | 19 | 13 | 14 | 15 | 16 |
| woodworking | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| zenotravel | 7 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 11 | 10 | 11 | 12 | 12 | 10 | 11 | 11 |
| Sum(1260) | 401 | 462 | 472 | 459 | 468 | 458 | 478 | 488 | 500 | 477 | 494 | 477 | 506 | 529 | 532 | 494 | 510 | 496 | 532 |

Table 1: Coverage (number of problems solved out of 1260) on Autoscale-21.11 IPC-based planning benchmark set (gripper, and miconic are replaced with harder instances. See S.5).

S.7 Comparisons Without Separate Generation and Evaluation (SGE)

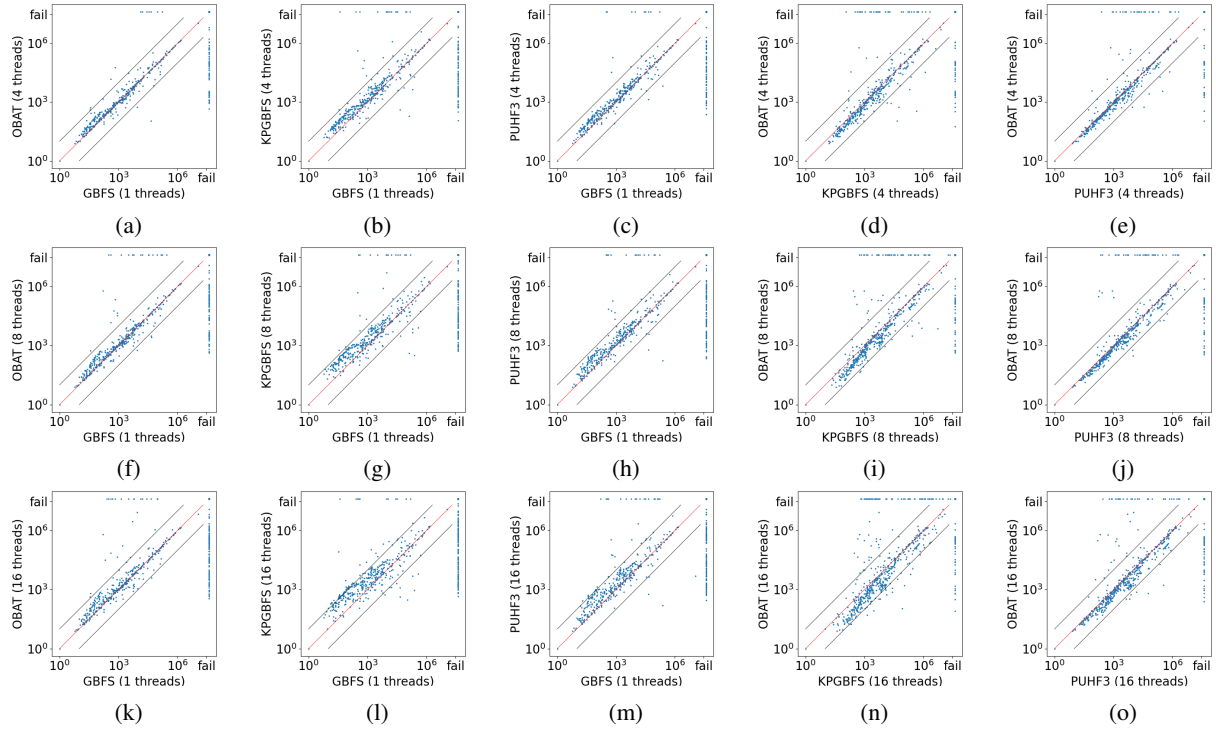


Figure 2: Number of states expanded, Diagonal lines are $y = 0.1x$, $y = x$, and $y = 10x$

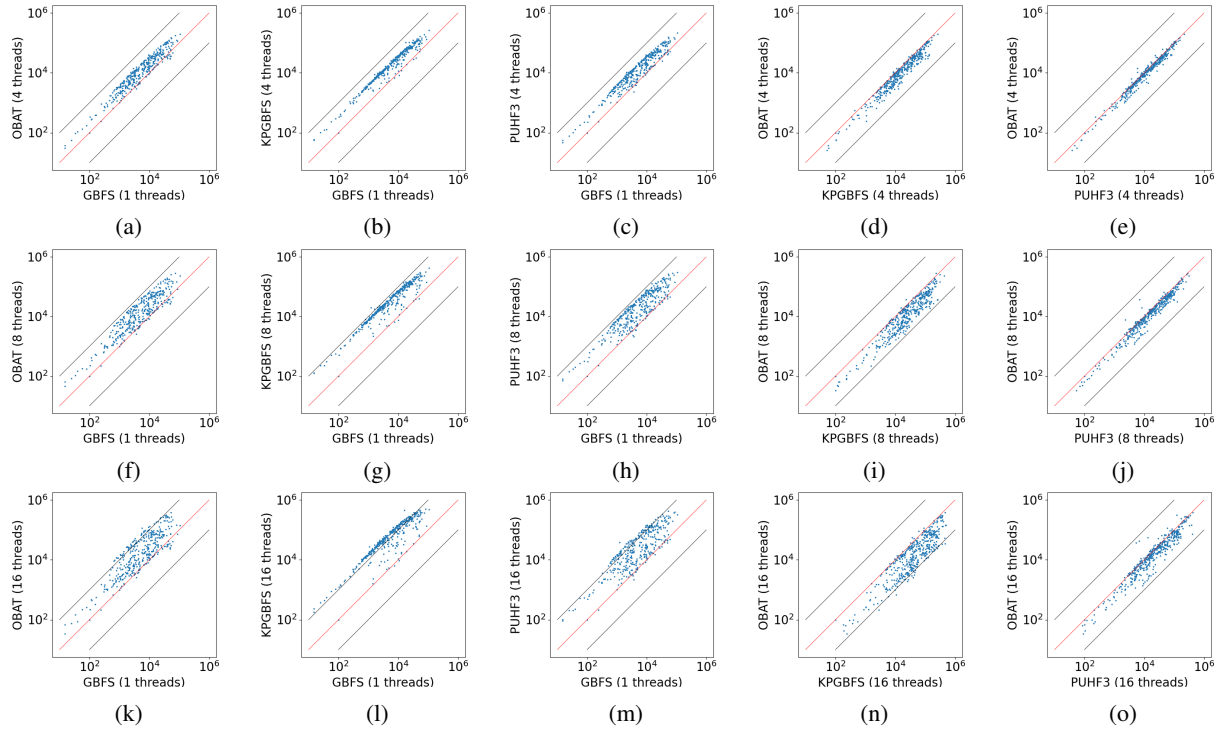


Figure 3: State evaluation rate comparison (states/second), Diagonal lines are $y = 0.1x$, $y = x$, and $y = 10x$

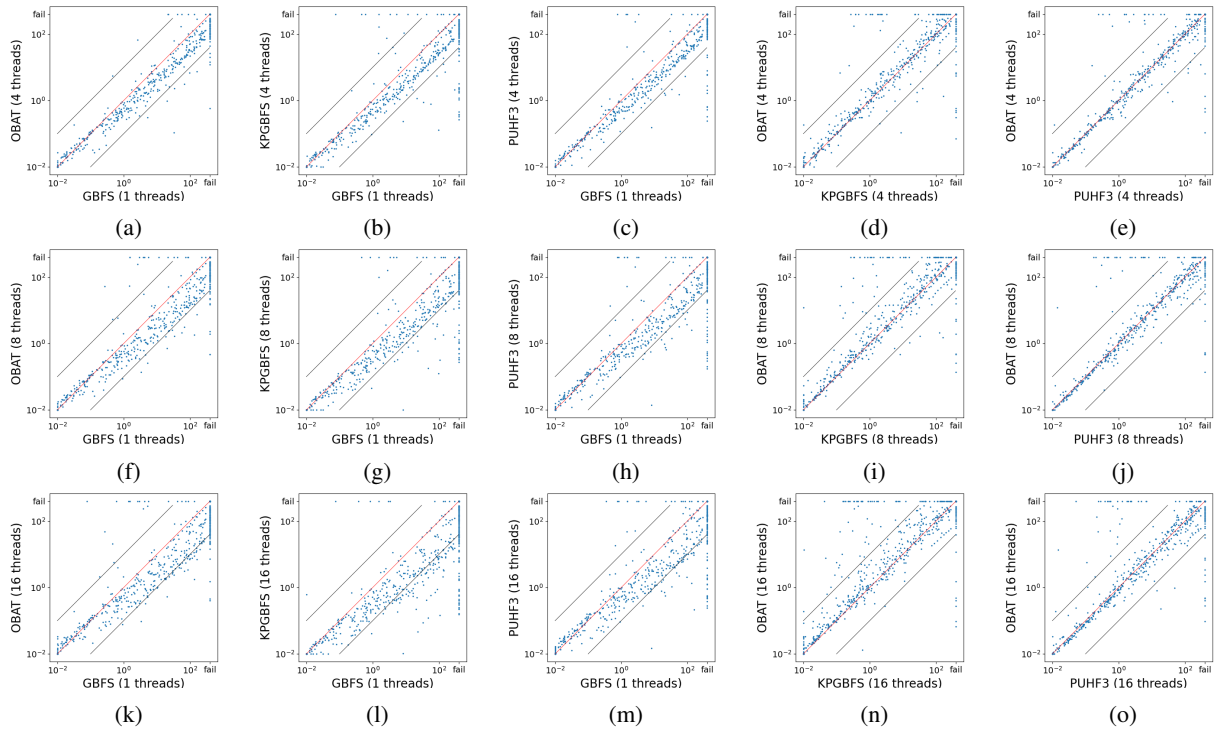


Figure 4: Search time (seconds) “fail”= out of time/memory, diagonal lines are $y = 0.1x$, $y = x$, and $y = 10x$

S.8 Comparisons Including Separate Generation and Evaluation (SGE)

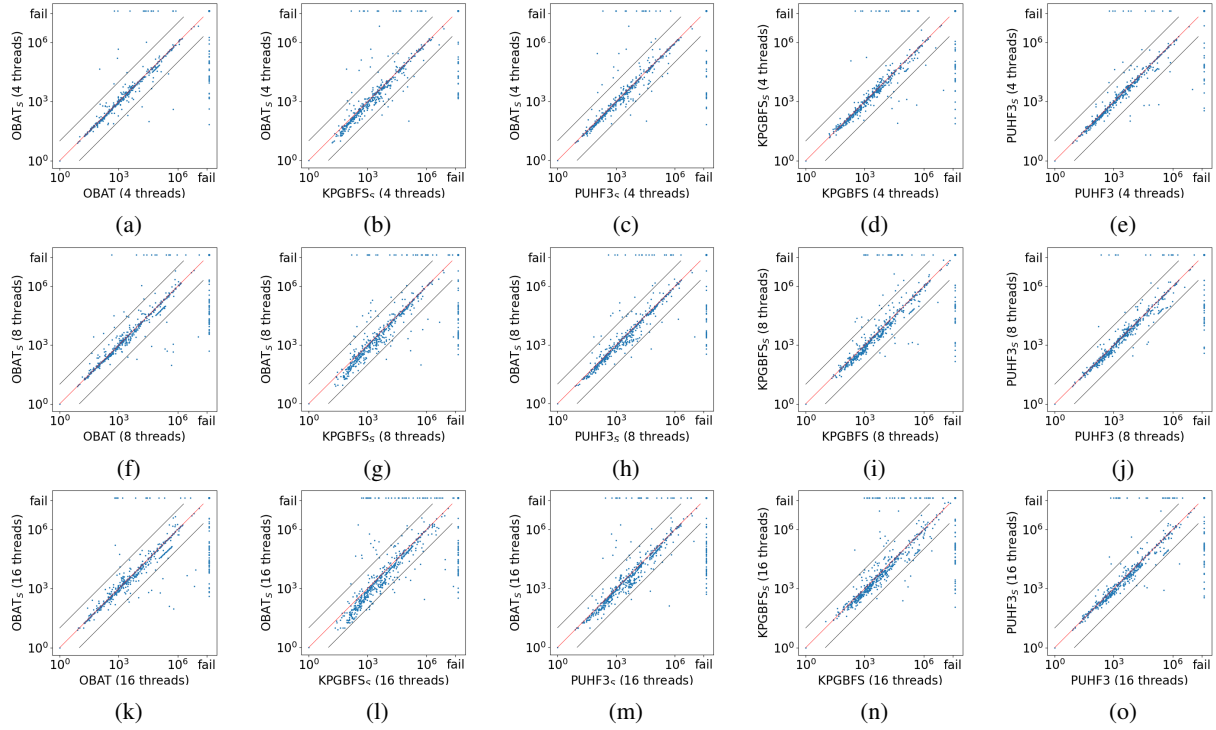


Figure 5: Number of states expanded, Diagonal lines are $y = 0.1x$, $y = x$, and $y = 10x$

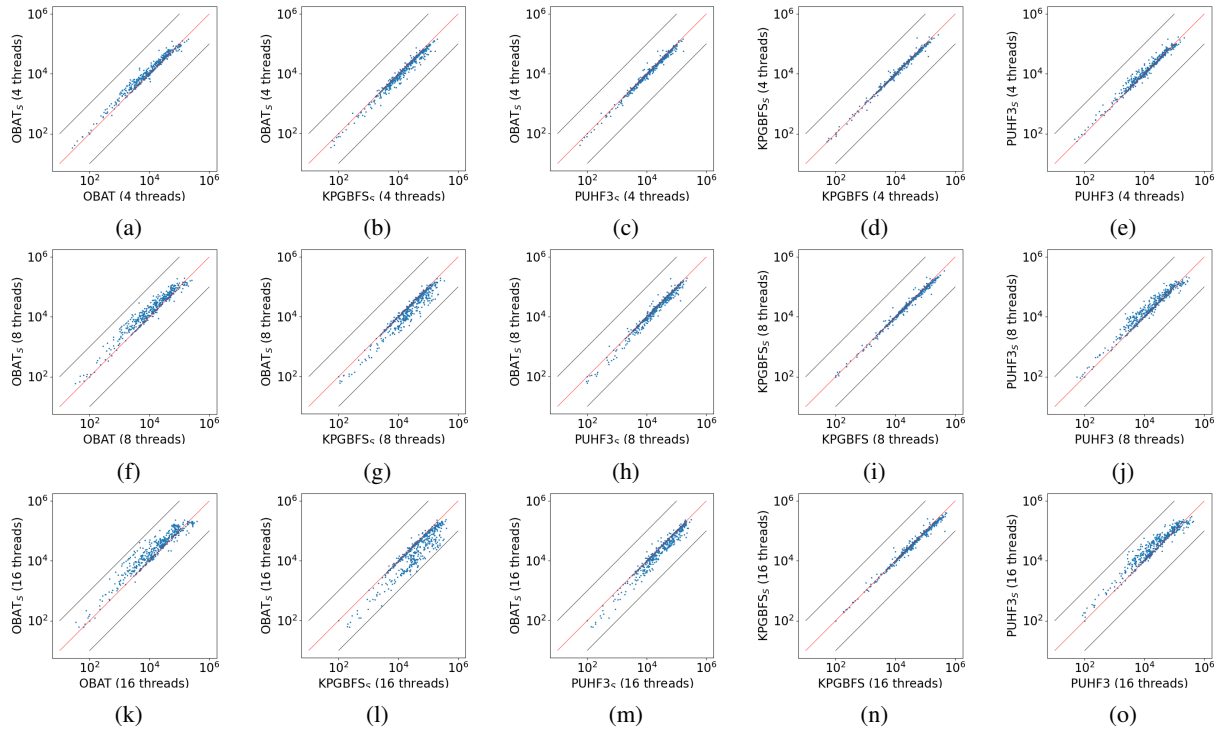


Figure 6: State evaluation rate comparison (states/second), Diagonal lines are $y = 0.1x$, $y = x$, and $y = 10x$

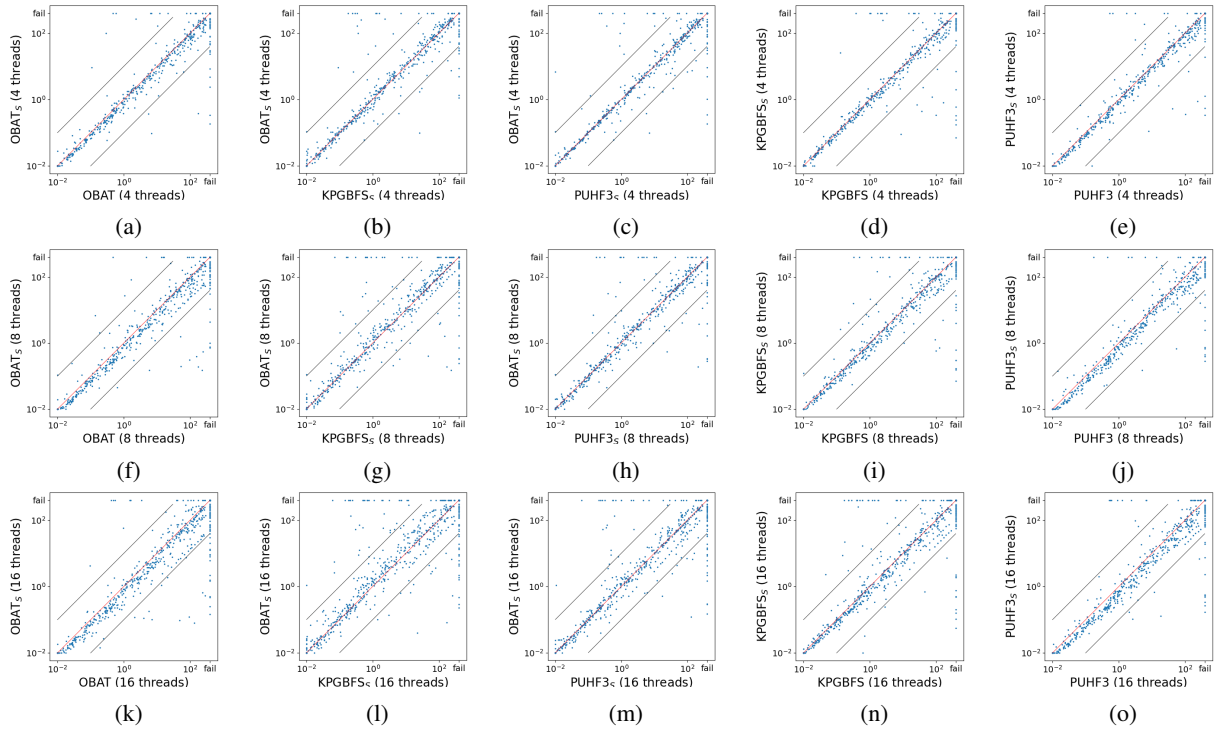


Figure 7: Search time (seconds) “fail”= out of time/memory, diagonal lines are $y = 0.1x$, $y = x$, and $y = 10x$

S.9 The performance impact of SGE on KPGBFS

105 Our experimental results show that KPGBFS_S (KPGBFS with SGE) had higher overall coverage (472/500/532 for 4/8/16 threads) than KPGBFS (462/488/529 for 4/8/16 threads).

In contrast, (Shimoda and Fukunaga 2024), the experimental evaluation of SGE on KPGBFS showed that KPGBFS_S had slightly lower total coverage (507/529/565 for 4/8/16 threads) than KPGBFS (510/534/567 for 4/8/16 threads).

110 These results are *not* contradictory – the difference is due to the fact that Shimoda and Fukunaga (2024) used the unmodified version of the AutoScale-21.11 benchmark set, while we used a modified version of the AutoScale-21.11 benchmark set. As explained in Section 7 and Supplement Section S.5, we replaced gripper and miconic with harder instances as all 30 of the original instances were solved by all methods and therefore unsuitable (too easy) for a performance evaluation.

115 If we remove the new gripper and miconic instances from the total coverage, and then replace them with the much easier, original, AutoScale-21.11 gripper and miconic instances used by (Shimoda and Fukunaga 2024), the total coverage for KPGBFS_S is 509/530/556 for 4/8/16 threads, and total coverage for KPGBFS is 509/532/573 for 4/8/16 threads, which is similar to the results in (Shimoda and Fukunaga 2024).